

# CSC 496 Mid-Project Report

**Learning to Play Snake: A Deep Q-Learning Study in a Custom Gym Environment**

**Gaurvendra Pratap Singh Pundhir**

## Abstract

I am training agents to play a custom 2-D Snake game using reinforcement learning. The environment is built from scratch with the Gymnasium API and supports both feature and pixel observations. The project proceeds from tabular baselines to Deep Q-Networks (DQN) with replay and target networks. By mid-term, the training/evaluation pipeline, logging, and checkpoints are complete. On the small-board configuration, DQN learns a stable policy; over the last 50 episodes of a representative run (seed 123) the **average return  $\approx 11.66$ , median survival  $\approx 136.0$  steps**, and **return standard deviation  $\approx 6.28$** . Next, I will add Double- and Dueling-DQN ablations, a modest curriculum, and stronger evaluation across multiple seeds.

---

## 1. Problem & Contributions

**Goal.** Learn policies that survive and collect apples efficiently in Snake, avoiding self-collision.

**How this project is different.**

- **Custom Gym environment** (feature and pixel modes) enabling both tabular and deep methods.
  - **Reproducible harness:** YAML configs, seeds, TensorBoard + CSV logging, and checkpoints.
  - **Systematic methodology:** tabular baselines → DQN → upcoming ablations (Double/Dueling/PER).
  - **Open, re-runnable code** suitable for peer replication and extension.
- 

## 2. Environment Formulation

**State space.**

- Feature state (current): compact features including head direction, relative food location, and local hazards; optionally frame-stacked.

- Pixel stack (optional): stack of the last k binary planes (snake/food/walls).

**Action space.** {Up, Down, Left, Right} with a guard against 180° reversals.

**Rewards.** +1 for apple, -1 for death, small per-step penalty (default -0.01, tuned to -0.003 for the best run). Optional potential-based shaping (progress toward food) is available without altering optimality.

**Termination.** Collision with boundary or self.

**API.** `reset()>(obs, info)`, `step(a)>(obs', r, terminated, truncated, info)` per Gymnasium.

This design satisfies MDP assumptions and scales from tabular methods to function approximation.

---

## 3. Methods

### 3.1 Baselines (v1)

Algorithms: Q-Learning / SARSA / Expected-SARSA (tabular).

Policy:  $\epsilon$ -greedy with linear decay.

Outcome: Learn simple survival on very small boards; quickly saturate due to state explosion and sparse reward.

### 3.2 Deep Q-Network (v2)

**Network (feature mode):** Flatten  $\rightarrow$  256 ReLU  $\rightarrow$  256 ReLU  $\rightarrow$   $|A|$  (MLP). (Pixel mode swaps MLP for a small CNN; not used in mid-term runs.)

**Stability:** Experience replay (uniform), target-network periodic sync, Huber loss, grad-norm clipping.

**Exploration:**  $\epsilon$  decays from 1.0  $\rightarrow$  0.05 over a chosen schedule.

**Typical hyperparameters:**  $\gamma = 0.99$ , batch = 128, buffer = 50–100k, warm-up = 8–10k, target sync = 1–2k, lr = 1e-3  $\rightarrow$  5e-4 (best stability), total steps = 150–300k.

**Planned ablations (next phase):** Double DQN (reduce overestimation), Dueling heads (value/advantage), Prioritized Replay (sample efficiency), and board-size curriculum ( $8 \times 8 \rightarrow 10 \times 10 \rightarrow 12 \times 12$ ).

---

## 4. Implementation & Reproducibility

**Repository:** `envs/` (custom Gym), `agents/` (DQN + Replay), `configs/` (YAMLs), `train_dqn.py`, `make_report.py`, `runs/` (TB + CSV), `checkpoints/`.

### Logging

- TensorBoard scalars: `charts/episode_return`, `charts/episode_length`, `loss/td_loss`.
- Per-episode CSV: `runs/episodes.csv` with `config`, `seed`, `episode`, `return`, `length`, `epsilon`, `loss`.
- Summary table: `make_report.py` aggregates the last K (default 50) episodes per (config, seed) into `runs/summary.csv`.

**Checkpoints:** `checkpoints/dqn_latest.pt` (periodic) and `dqn_final.pt` (end of run).

**Seeding:** NumPy and PyTorch seeds set from CLI (`--seed`).

## Quickstart (PowerShell)

- `conda create -n RL python=3.11 -y`
  - `conda activate RL`
  - `python -m pip install -r requirements.txt`
  - `python train_dqn.py --config configs\dqn_small.yaml --seed 123`
  - `tensorboard --logdir runs`
  - `python make_report.py`
- 

## 5. Experiments

### 5.1 Configurations

**Small board (feature state)** — best-performing mid-term setup.

Key values that improved stability: `step_penalty: -0.003`, `lr: 0.0005`, `target_sync: 2000`, `eps_decay_steps: 150000`.

### 5.2 Evaluation Protocol

Train with seed(s) per config; log every episode to CSV and TensorBoard.

Report metrics from the last 50 episodes per (config, seed): **average return** ( $\uparrow$ ), **median survival steps** ( $\uparrow$ ), and **return std** ( $\downarrow$ ).

Qualitatively inspect learning curves (EMA-smoothed plots) for upward trends and stability.

---

## 6. Results (Mid-Term)

### 6.1 Quantitative (seed 123, last-50 episodes)

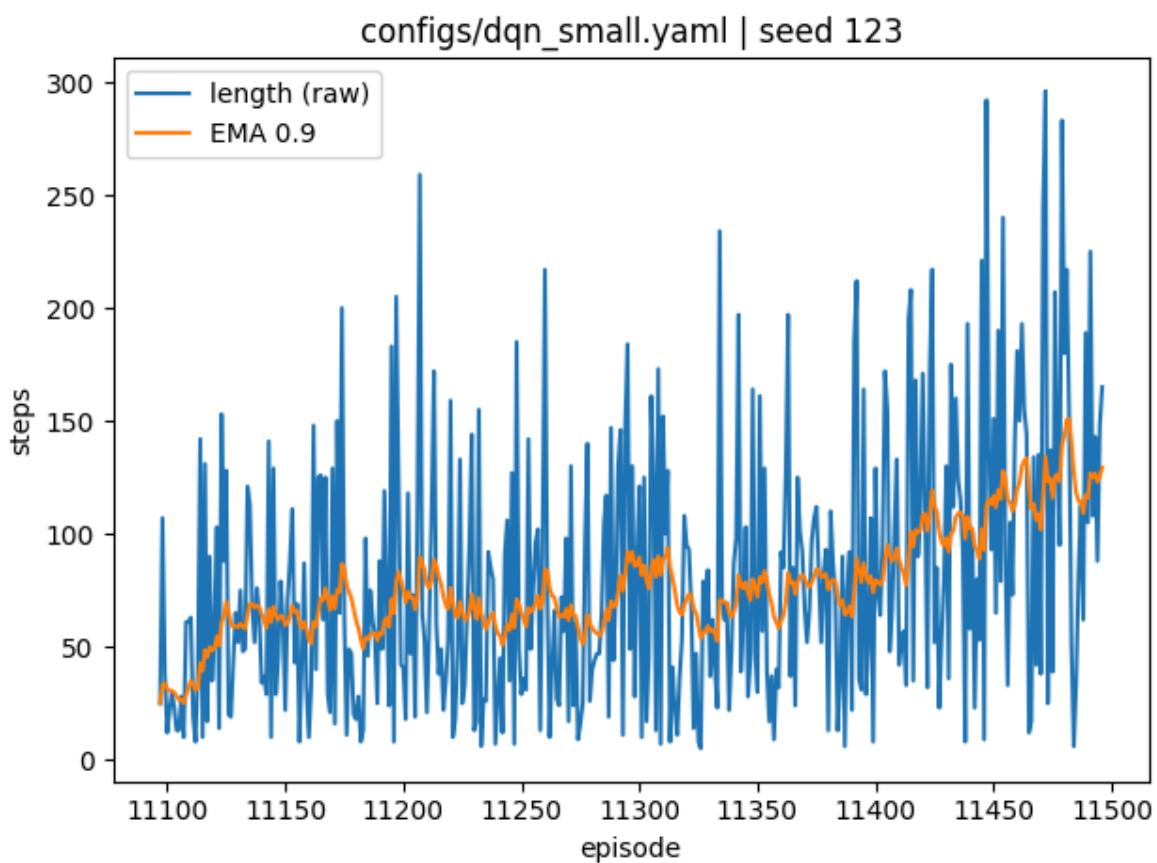
Config	S	Avg. Return	Median Survival	Return Std
configs/dqn_small.yaml ml	1	<b>11.6628</b>	<b>136.0</b>	<b>6.2820</b>

(Derived from `runs/summary.csv` after running `make_report.py` on the current `runs/episodes.csv`.)

## 6.2 Qualitative

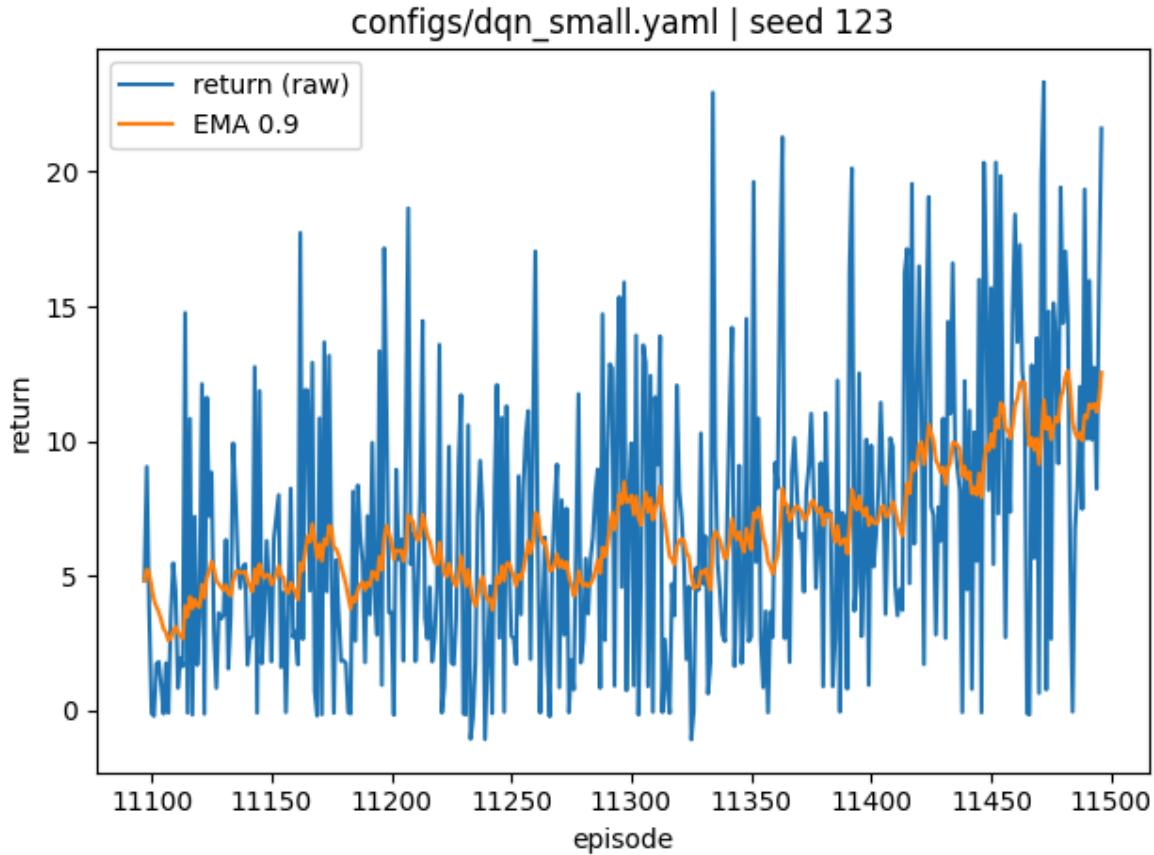
**Figure 1 – Episode Return (last ~400 episodes, EMA 0.9).**

Shows a clear upward trend; the smoothed curve rises steadily across training.



**Figure 2 – Episode Length (last ~400 episodes, EMA 0.9).**

Median survival increases with frequent long episodes (>100–200 steps) by the end.



---

## 7. Discussion

**Learning achieved.** DQN transitions from survival-only behavior to purposeful apple collection on the small board.

**What helped.** Gentler step penalty, slower  $\epsilon$ -decay, slightly lower lr, and longer target-sync interval reduced instability and prevented “wander with cost” failure modes.

**Variance.** Return variance remains non-trivial ( $\text{std} \approx 6.28$ ) due to stochastic starts and sparse reward; averaging across seeds and/or Double DQN should tighten dispersion.

**Limitations.** Current results use feature state; pixel stacks and generalization to larger boards are pending. Baselines were used for calibration but not extensively profiled in mid-term plots.

---

## 8. Risks & Mitigations

Risk	Impact	Mitigation

Overestimation bias	Noisy Q-values, unstable returns	<b>Double DQN</b> target selection
Sparse rewards	Long plateaus, looping behavior	Potential-based shaping; curriculum; tuned step penalty
Seed variance	Unreliable single-run conclusions	Report across $\geq 3$ seeds; fixed eval seeds
Scaling to larger boards	Slower learning	Curriculum ( $8 \times 8 \rightarrow 10 \times 10 \rightarrow 12 \times 12$ ); PER; dueling heads

---

## 9. Next-Phase Plan (2–3 weeks)

**Algorithms.** Implement **Double DQN**; add **Dueling** head; consider **Prioritized Replay** with  $\beta$ -annealing.

**Training.** Curriculum  $8 \times 8 \rightarrow 10 \times 10 \rightarrow 12 \times 12$ ; grid search (step penalty,  $\epsilon$  schedule, lr, target\_sync).

**Reporting.** Standardized evaluation ( $\geq 3$  seeds), learning curves per config, ablation tables, demo GIF.

---

## 10. Timeline (updated)

Date	Milestone	Deliverables
Oct 17	Proposal & Literature Review	Proposal doc + references
<b>Nov 10</b>	<b>Mid-Project</b>	Working env + DQN results + this report + plots

Nov 24	Ablations Complete	Double/Dueling/PER + curriculum; updated tables/plots
Dec 5	Final Report	6–8 pp paper + code + demo GIF/video

---

## 11. References

Mnih, V. et al. Human-level control through deep reinforcement learning. **Nature**, 2015.

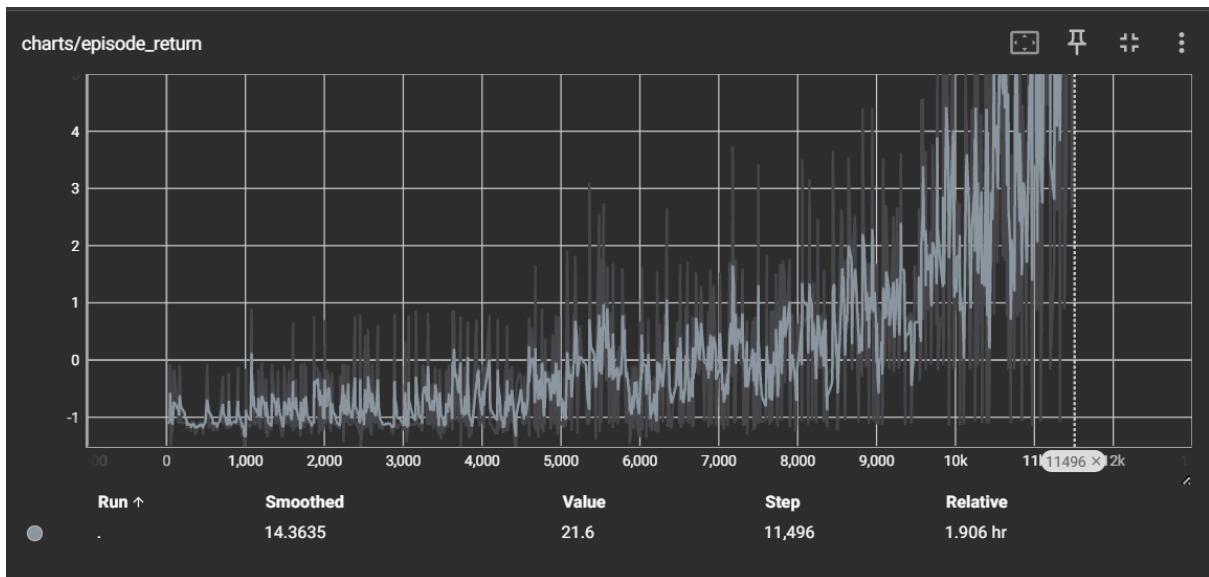
Gymnasium documentation — <https://gymnasium.farama.org> (API reference).

Hessel, M. et al. Rainbow: Combining Improvements in Deep RL. **AAAI**, 2018.

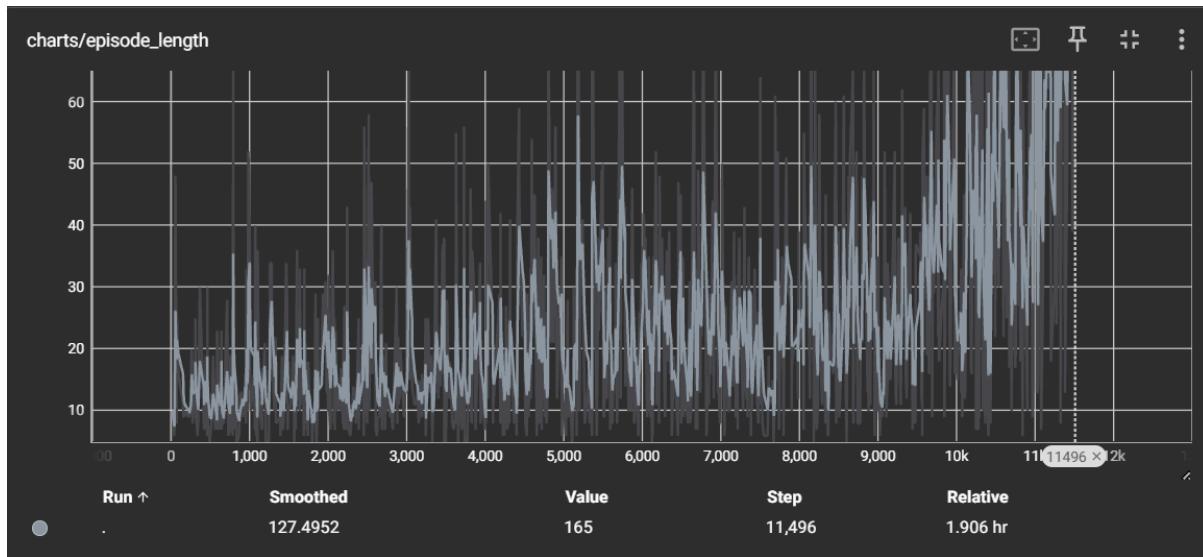
---

## 12. Additional Figures (TensorBoard, seed 123)

**Figure 3 — Episode Return (TB export).** Return rises from near 0 to >4–5 by ~11.5k episodes; EMA shows steady improvement.



**Figure 4 — Episode Length (TB export).** Survival increases with frequent long episodes (50–150+ steps), consistent with improved navigation.

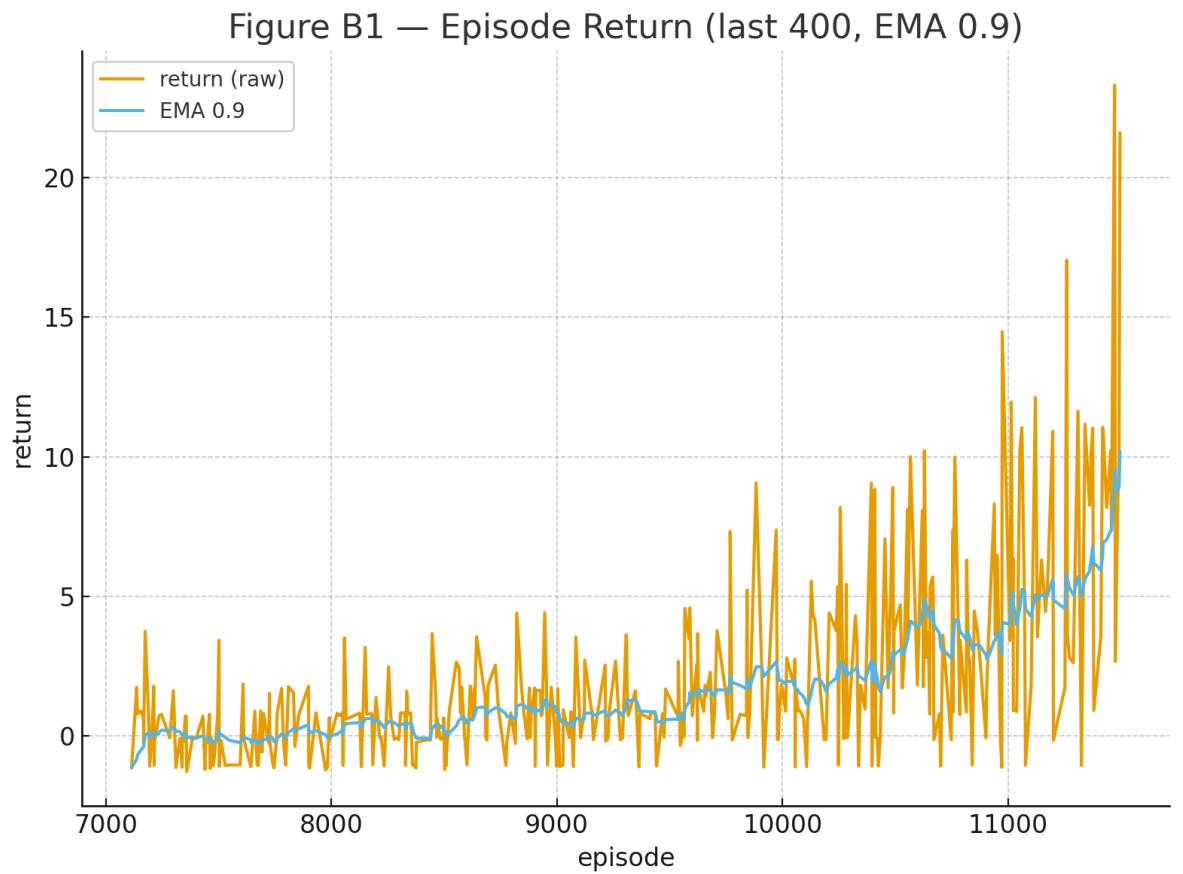


## Appendix A — Commands Used

- # Train (small board)
- python train\_dqn.py --config configs\dqn\_small.yaml --seed 123
- 
- # Live metrics
- tensorboard --logdir runs
- 
- # Build table & plots
- python make\_report.py
- python make\_plots.py --csv runs\episodes.csv --outdir submission --last 400 --smooth 0.9
- 
- # Collect artifacts
- Copy-Item runs\summary.csv submission\
- Copy-Item runs\episodes.csv submission\episodes\_full.csv
- Copy-Item configs\dqn\_small.yaml submission\config\_used.yaml
- Copy-Item checkpoints\dqn\_final.pt submission\ -ErrorAction SilentlyContinue
- Copy-Item checkpoints\dqn\_latest.pt submission\ -ErrorAction SilentlyContinue
- Compress-Archive -Path submission\\* -DestinationPath submission\_midproject.zip -Force

## Appendix B — Raw Plots

- **Figure B1 — Episode Return (last 400, EMA 0.9).**



- **Figure B2 — Episode Length (last 400, EMA 0.9).**

