

CSC 496 Mid-Project Report

Learning to Play Snake: A Deep Q-Learning Study in a Custom Gym Environment

Gaurvendra Pratap Singh Pundhir

Abstract

I am training agents to play a custom 2-D Snake game using reinforcement learning. The environment is built from scratch with the Gymnasium API and supports both feature and pixel observations. The project proceeds from tabular baselines to Deep Q-Networks (DQN) with replay and target networks. By mid-term, the training/evaluation pipeline, logging, and checkpoints are complete. On the small-board configuration, DQN learns a stable policy; over the last 50 episodes of a representative run (seed 123) the **average return ≈ 11.66 , median survival ≈ 136.0 steps**, and **return standard deviation ≈ 6.28** . Next, I will add Double- and Dueling-DQN ablations, a modest curriculum, and stronger evaluation across multiple seeds.

1. Problem & Contributions

Goal. Learn policies that survive and collect apples efficiently in Snake, avoiding self-collision.

How this project is different.

- **Custom Gym environment** (feature and pixel modes) enabling both tabular and deep methods.
 - **Reproducible harness:** YAML configs, seeds, TensorBoard + CSV logging, and checkpoints.
 - **Systematic methodology:** tabular baselines → DQN → upcoming ablations (Double/Dueling/PER).
 - **Open, re-runnable code** suitable for peer replication and extension.
-

2. Environment Formulation

State space.

- Feature state (current): compact features including head direction, relative food location, and local hazards; optionally frame-stacked.

- Pixel stack (optional): stack of the last k binary planes (snake/food/walls).

Action space. {Up, Down, Left, Right} with a guard against 180° reversals.

Rewards. +1 for apple, -1 for death, small per-step penalty (default -0.01, tuned to -0.003 for the best run). Optional potential-based shaping (progress toward food) is available without altering optimality.

Termination. Collision with boundary or self.

API. `reset()>(obs, info)`, `step(a)>(obs', r, terminated, truncated, info)` per Gymnasium.

This design satisfies MDP assumptions and scales from tabular methods to function approximation.

3. Methods

3.1 Baselines (v1)

Algorithms: Q-Learning / SARSA / Expected-SARSA (tabular).

Policy: ϵ -greedy with linear decay.

Outcome: Learn simple survival on very small boards; quickly saturate due to state explosion and sparse reward.

3.2 Deep Q-Network (v2)

Network (feature mode): Flatten \rightarrow 256 ReLU \rightarrow 256 ReLU \rightarrow $|A|$ (MLP). (Pixel mode swaps MLP for a small CNN; not used in mid-term runs.)

Stability: Experience replay (uniform), target-network periodic sync, Huber loss, grad-norm clipping.

Exploration: ϵ decays from 1.0 \rightarrow 0.05 over a chosen schedule.

Typical hyperparameters: $\gamma = 0.99$, batch = 128, buffer = 50–100k, warm-up = 8–10k, target sync = 1–2k, lr = 1e-3 \rightarrow 5e-4 (best stability), total steps = 150–300k.

Planned ablations (next phase): Double DQN (reduce overestimation), Dueling heads (value/advantage), Prioritized Replay (sample efficiency), and board-size curriculum ($8 \times 8 \rightarrow 10 \times 10 \rightarrow 12 \times 12$).

4. Implementation & Reproducibility

Repository: `envs/` (custom Gym), `agents/` (DQN + Replay), `configs/` (YAMLs), `train_dqn.py`, `make_report.py`, `runs/` (TB + CSV), `checkpoints/`.

Logging

- TensorBoard scalars: `charts/episode_return`, `charts/episode_length`, `loss/td_loss`.
- Per-episode CSV: `runs/episodes.csv` with `config`, `seed`, `episode`, `return`, `length`, `epsilon`, `loss`.
- Summary table: `make_report.py` aggregates the last K (default 50) episodes per (config, seed) into `runs/summary.csv`.

Checkpoints: `checkpoints/dqn_latest.pt` (periodic) and `dqn_final.pt` (end of run).

Seeding: NumPy and PyTorch seeds set from CLI (`--seed`).

Quickstart (PowerShell)

- `conda create -n RL python=3.11 -y`
 - `conda activate RL`
 - `python -m pip install -r requirements.txt`
 - `python train_dqn.py --config configs\dqn_small.yaml --seed 123`
 - `tensorboard --logdir runs`
 - `python make_report.py`
-

5. Experiments

5.1 Configurations

Small board (feature state) — best-performing mid-term setup.

Key values that improved stability: `step_penalty: -0.003`, `lr: 0.0005`, `target_sync: 2000`, `eps_decay_steps: 150000`.

5.2 Evaluation Protocol

Train with seed(s) per config; log every episode to CSV and TensorBoard.

Report metrics from the last 50 episodes per (config, seed): **average return** (\uparrow), **median survival steps** (\uparrow), and **return std** (\downarrow).

Qualitatively inspect learning curves (EMA-smoothed plots) for upward trends and stability.

6. Results (Mid-Term)

6.1 Quantitative (seed 123, last-50 episodes)

Config	S	Avg. Return	Median Survival	Return Std
configs/dqn_small.yaml ml	1	11.6628	136.0	6.2820

(Derived from `runs/summary.csv` after running `make_report.py` on the current `runs/episodes.csv`.)

6.2 Qualitative

Figure 1 – Episode Return (last ~400 episodes, EMA 0.9).

Shows a clear upward trend; the smoothed curve rises steadily across training.

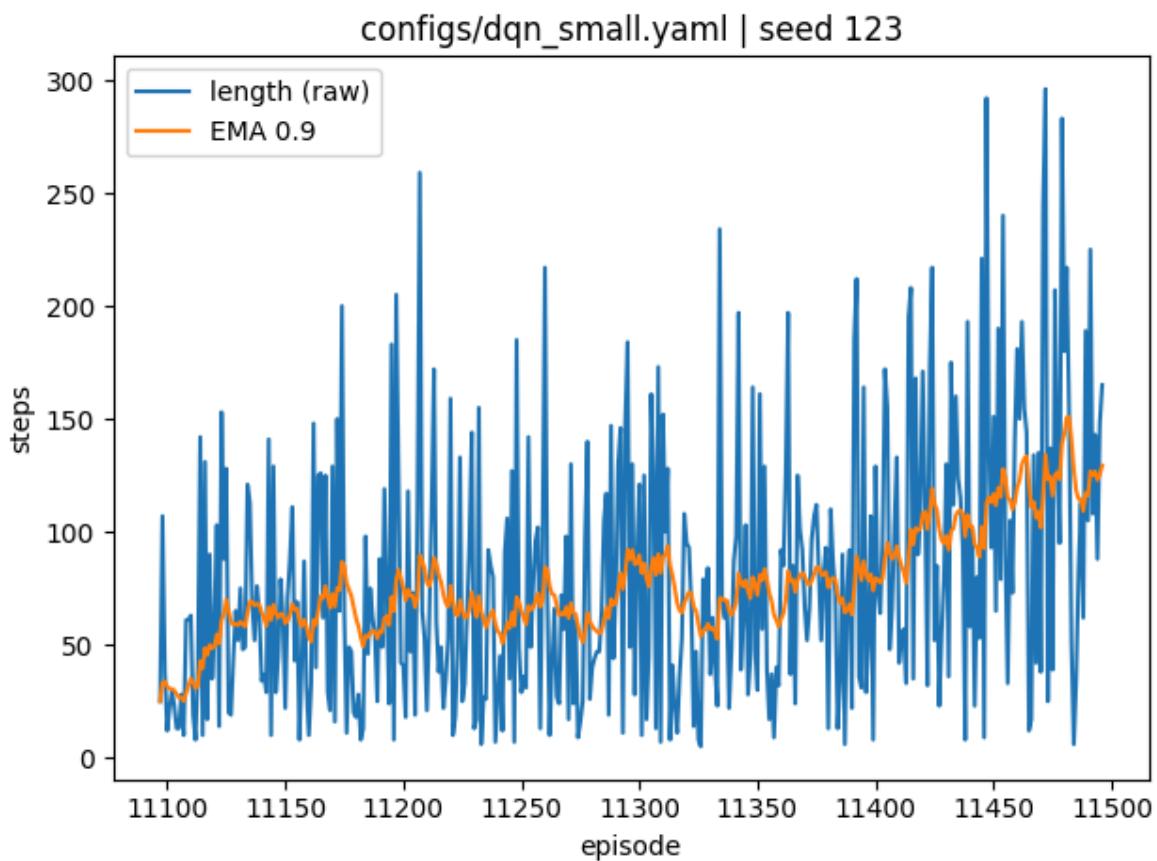
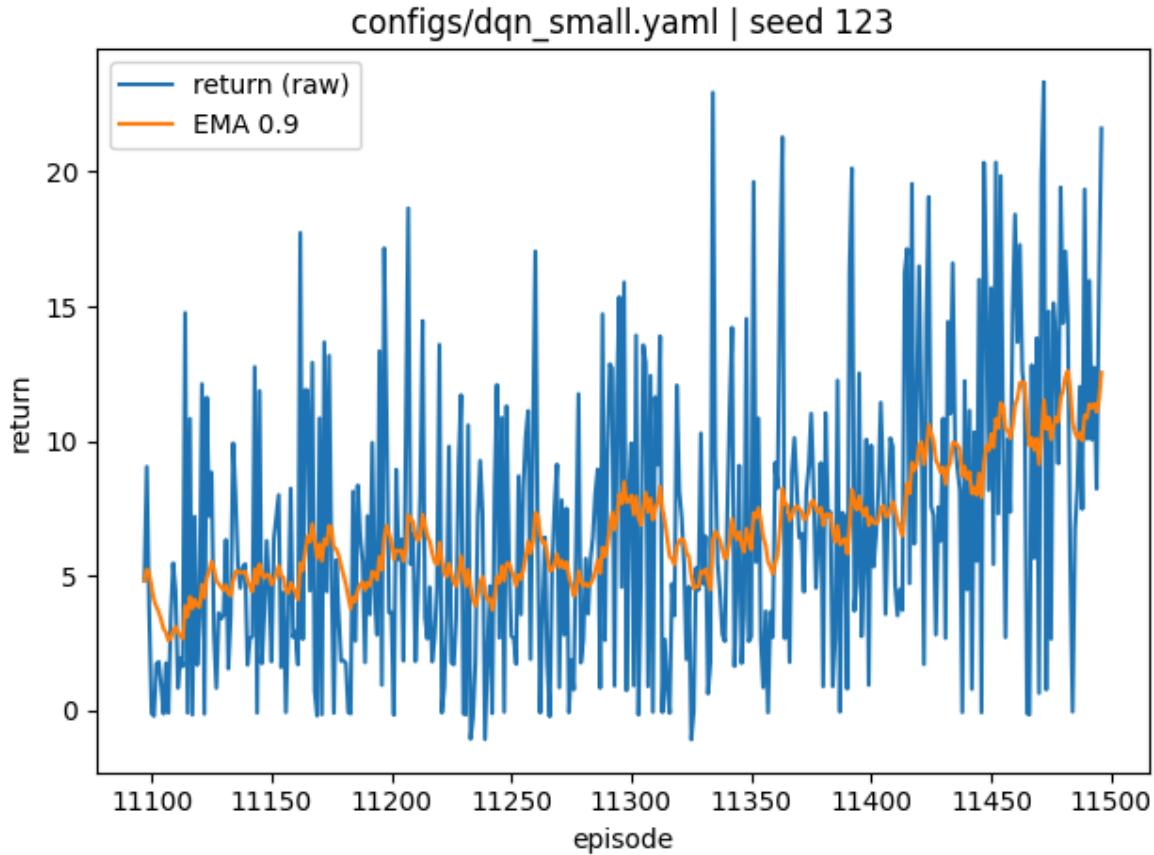


Figure 2 – Episode Length (last ~400 episodes, EMA 0.9).

Median survival increases with frequent long episodes (>100–200 steps) by the end.



7. Discussion

Learning achieved. DQN transitions from survival-only behavior to purposeful apple collection on the small board.

What helped. Gentler step penalty, slower ϵ -decay, slightly lower lr, and longer target-sync interval reduced instability and prevented “wander with cost” failure modes.

Variance. Return variance remains non-trivial ($\text{std} \approx 6.28$) due to stochastic starts and sparse reward; averaging across seeds and/or Double DQN should tighten dispersion.

Limitations. Current results use feature state; pixel stacks and generalization to larger boards are pending. Baselines were used for calibration but not extensively profiled in mid-term plots.

8. Risks & Mitigations

Risk

Impact

Mitigation

Overestimation bias	Noisy Q-values, unstable returns	Double DQN target selection
Sparse rewards	Long plateaus, looping behavior	Potential-based shaping; curriculum; tuned step penalty
Seed variance	Unreliable single-run conclusions	Report across ≥ 3 seeds; fixed eval seeds
Scaling to larger boards	Slower learning	Curriculum ($8 \times 8 \rightarrow 10 \times 10 \rightarrow 12 \times 12$); PER; dueling heads

9. Next-Phase Plan (2–3 weeks)

Algorithms. Implement **Double DQN**; add **Dueling** head; consider **Prioritized Replay** with β -annealing.

Training. Curriculum $8 \times 8 \rightarrow 10 \times 10 \rightarrow 12 \times 12$; grid search (step penalty, ϵ schedule, lr, target_sync).

Reporting. Standardized evaluation (≥ 3 seeds), learning curves per config, ablation tables, demo GIF.

10. Timeline (updated)

Date	Milestone	Deliverables
Oct 17	Proposal & Literature Review	Proposal doc + references
Nov 10	Mid-Project	Working env + DQN results + this report + plots

Nov 24	Ablations Complete	Double/Dueling/PER + curriculum; updated tables/plots
Dec 5	Final Report	6–8 pp paper + code + demo GIF/video

11. References

Mnih, V. et al. Human-level control through deep reinforcement learning. **Nature**, 2015.

Gymnasium documentation — <https://gymnasium.farama.org> (API reference).

Hessel, M. et al. Rainbow: Combining Improvements in Deep RL. **AAAI**, 2018.

12. Additional Figures (TensorBoard, seed 123)

Figure 3 — Episode Return (TB export). Return rises from near 0 to >4–5 by ~11.5k episodes; EMA shows steady improvement.

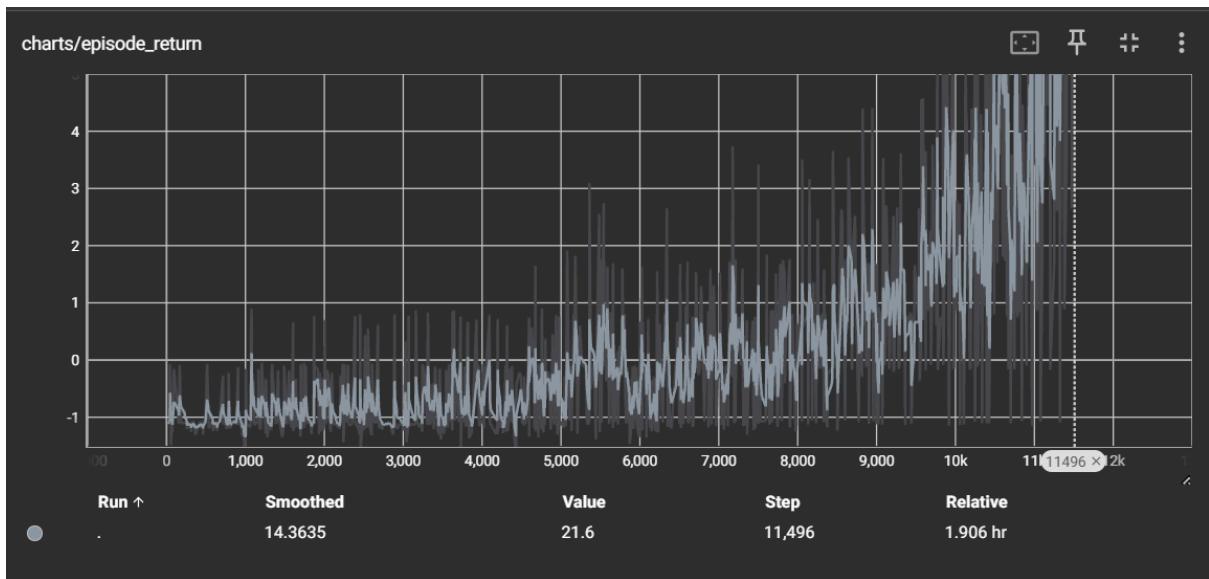
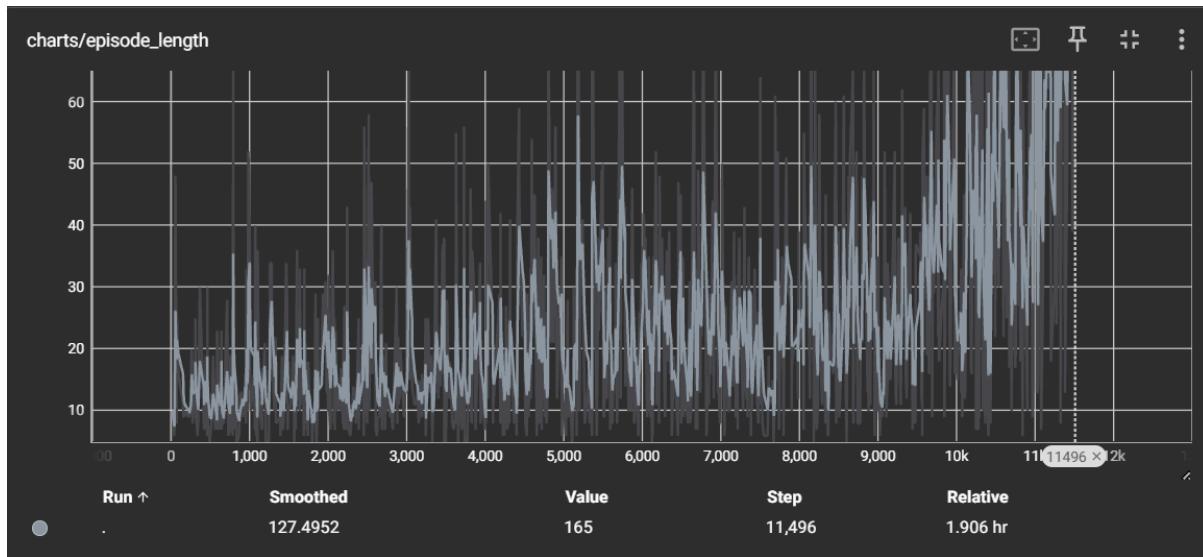


Figure 4 — Episode Length (TB export). Survival increases with frequent long episodes (50–150+ steps), consistent with improved navigation.

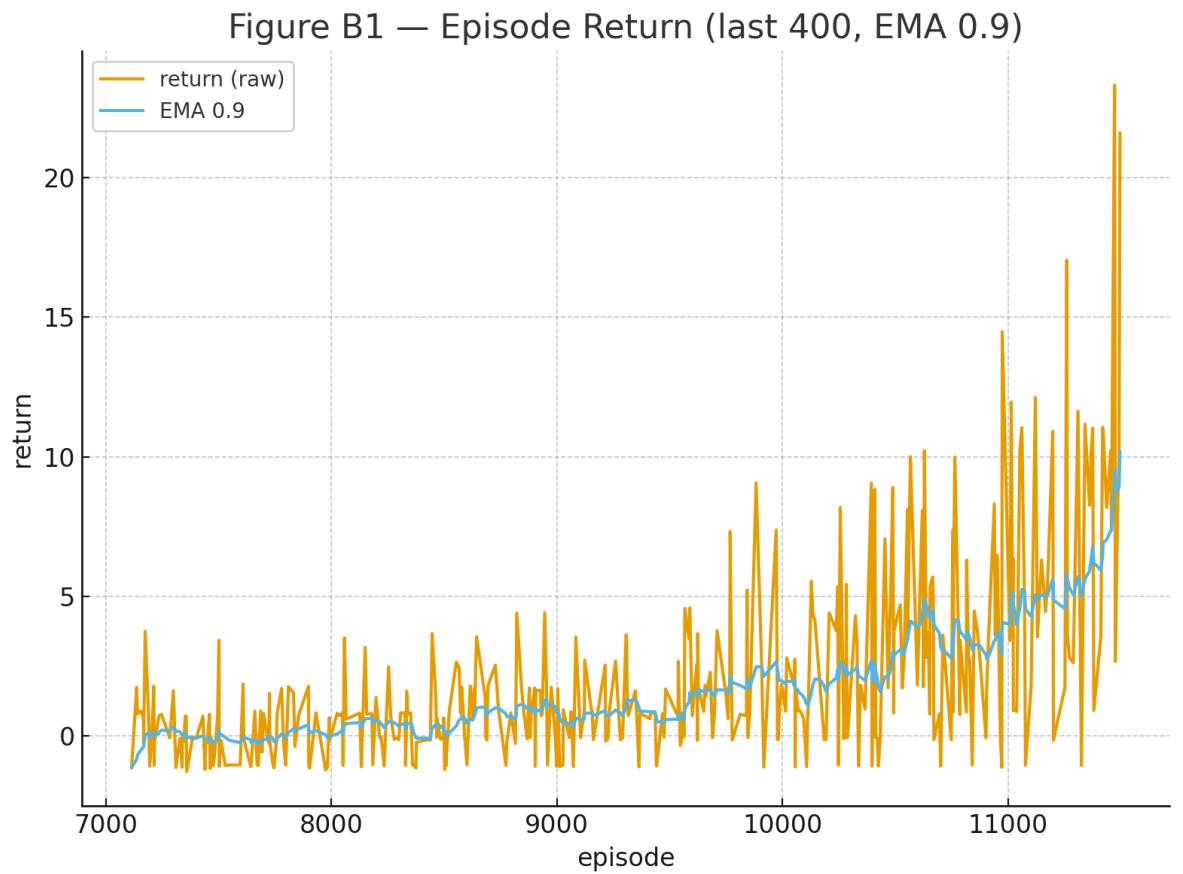


Appendix A — Commands Used

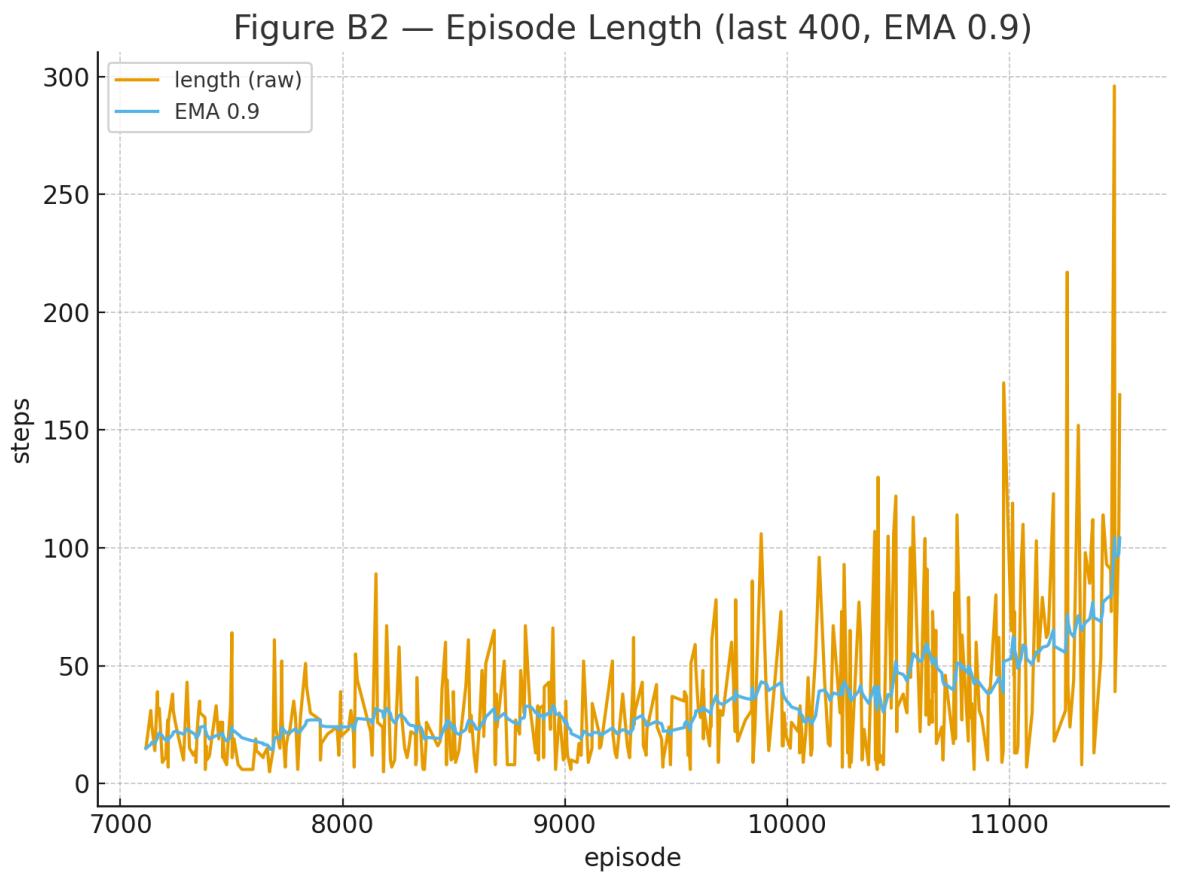
- # Train (small board)
- python train_dqn.py --config configs\dqn_small.yaml --seed 123
-
- # Live metrics
- tensorboard --logdir runs
-
- # Build table & plots
- python make_report.py
- python make_plots.py --csv runs\episodes.csv --outdir submission --last 400 --smooth 0.9
-
- # Collect artifacts
- Copy-Item runs\summary.csv submission\
- Copy-Item runs\episodes.csv submission\episodes_full.csv
- Copy-Item configs\dqn_small.yaml submission\config_used.yaml
- Copy-Item checkpoints\dqn_final.pt submission\ -ErrorAction SilentlyContinue
- Copy-Item checkpoints\dqn_latest.pt submission\ -ErrorAction SilentlyContinue
- Compress-Archive -Path submission* -DestinationPath submission_midproject.zip -Force

Appendix B — Raw Plots

- **Figure B1 — Episode Return (last 400, EMA 0.9).**



- **Figure B2 — Episode Length (last 400, EMA 0.9).**



Project Snapshot

The screenshot shows a terminal window with the following command and its output:

```
snake_rl_final_ready
```

```
summary.csv runs plot_return_seed123.png episodes_full.csv summary.csv submission
```

```
submission > episodes_full.csv > data
```

```
1 config,seed,episode,return,length,epsilon,loss
2 configs/dqn_small.yaml,123,0,-1.09,10,0.9999715,0.0
3 configs/dqn_small.yaml,123,1,-1.2,21,0.999985,0.0
4 configs/dqn_small.yaml,123,2,-1.08,9,0.9998765000000001,0.0
5 configs/dqn_small.yaml,123,3,-1.15,16,0.9998258333333333,0.0
```

The terminal also displays a large amount of training log output, which is mostly illegible due to the high character count. It includes metrics like training percentage, time, and loss values.

- There isn't a git commit history - I haven't put it on github yet as I'm the sole collaborator. So, I didn't had any need to branch the project out - for the final submission - I'll put the files on github, and share the link